

# Model\_Training

August 31, 2020

```
[22]: import numpy as np
import pandas as pd
import sklearn
import scipy
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.metrics import classification_report, accuracy_score
from sklearn.ensemble import IsolationForest
from sklearn.neighbors import LocalOutlierFactor
from sklearn.svm import OneClassSVM

[2]: # Read the CSV data into the pandas dataframe
file = '../data/data.csv'
data = pd.read_csv(file)

[3]: # Let's replace the label column values to 0 for normal and 1 for others!

LABELS = data["label"].unique() # Get Unique Values of label column
LABELS = [label for label in LABELS if label != "normal"] #All labels other
↳ than Normal!

data["label"].replace(['normal'], 0, inplace=True)
data["label"].replace(LABELS, 1, inplace=True)

[4]: data = data.fillna("None") # fill all the empty hosts with 0

#Now let's encode all the columns which are not float or int e.g. sourceIP, MAC
↳ to make it possible for model to interpret
columnsToEncode = list(data.select_dtypes(include=['category', 'object']))

le = LabelEncoder() # use label encoder from sklearn

for feature in columnsToEncode:
    try:
        data[feature] = le.fit_transform(data[feature])
```

```

        #print(data[feature])
    except:
        print ('error' + feature)

```

```
[5]: data.head()
```

```

[5]:   uid  sourceMac  sourceIp  destIp  destMac  sourcePort  destPort  host  kIn  \
0    0          0          2     15         5        19357        8000    3  0.0
1    0          0          2     15         5         1939        8000    3  0.0
2    0          0          2     15         5        19668        8000    3  0.0
3    0          0          2     15         5        19807        8000    3  0.0
4    0          0          2     15         5        19851        8000    3  0.0

      kOut  ...  outPacketsNo  protocol  urgent  ack  push  reset  syn  fin  \
0  1.015625  ...             2         17        0    0    0        0    0    0
1  1.015625  ...             2         17        0    0    0        0    0    0
2  0.507813  ...             1         17        0    0    0        0    0    0
3  0.507813  ...             1         17        0    0    0        0    0    0
4  0.507813  ...             1         17        0    0    0        0    0    0

      timestamp  label
0         22374      1
1         22374      1
2         22374      1
3         22374      1
4         22374      1

```

[5 rows x 21 columns]

```
[6]: data["label"].value_counts()
```

```

[6]: 1    2988506
     0     91386
     Name: label, dtype: int64

```

```

[7]: Train, Val = sklearn.model_selection.train_test_split(data, test_size=0.9,
↳ random_state=1, shuffle=True)

```

```
[8]: Train["label"].value_counts()
```

```

[8]: 1    298792
     0     9197
     Name: label, dtype: int64

```

```
[9]: Val["label"].value_counts()
```

```
[9]: 1    2689714
      0     82189
      Name: label, dtype: int64
```

```
[10]: Attack = Train[Train['label']==1]
      Normal = Train[Train['label']==0]
      outlier_fraction = np.ceil(len(Attack)/float(len(Normal)))

      #Let's print how many more outliers are there in the dataset compared to normal
      ↪data
      print(outlier_fraction)
```

33.0

```
[35]: #Let's now split the features and the target ground truth
      features = [feature for feature in Train.columns.tolist() if feature not in
      ↪["label"]]
      target = "label"

      # Define a random state
      state = np.random.RandomState(42)
      X_train = Train[features]
      Y_train = Train[target]

      X_val = Val[features]
      Y_val = Val[target]

      X_outliers = state.uniform(low=0, high=1, size=(X_train.shape[0], X_train.
      ↪shape[1]))
      # Print the shapes of X & Y
      print(X_train.shape)
      print(Y_train.shape)
```

(307989, 20)

(307989,)

```
[13]: #Let's initialize some classifiers
      classifiers = {
          "Isolation Forest":IsolationForest(n_estimators=100,
      ↪max_samples=len(X_train),
          ↪
          ↪contamination=outlier_fraction,random_state=state, verbose=1),
          "Local Outlier Factor":LocalOutlierFactor(n_neighbors=20, algorithm='auto',
          leaf_size=30, metric='minkowski',
          p=2, metric_params=None,
      ↪contamination=0.03),
```

```

        "Support Vector Machine":OneClassSVM(kernel='rbf', degree=3, gamma=0.1,nu=0.
↪05,

                                         max_iter=-1 )

}

```

```

[ ]: n_outliers = len(Attack)
for i, (clf_name,clf) in enumerate(classifiers.items()):
    #Fit the data and tag outliers
    if clf_name == "Local Outlier Factor":
        continue
    y_pred = clf.fit_predict(X_train)
    scores_prediction = clf.negative_outlier_factor_
    elif clf_name == "Support Vector Machine":
        clf.fit(X_train)
        y_pred = clf.predict(X_train)
    else:
        continue
    clf.fit(X_train)
    scores_prediction = clf.decision_function(X_train)
    y_pred = clf.predict(X_train)

    #Reshape the prediction values to 0 for Valid transactions , 1 for Fraud
↪transactions
    y_pred[y_pred == 1] = 0
    y_pred[y_pred == -1] = 1
    n_errors = (y_pred != Y_train).sum()
    # Run Classification Metrics
    print("{}: {}".format(clf_name,n_errors))
    print("Accuracy Score :")
    print(accuracy_score(Y_train,y_pred))
    print("Classification Report :")
    print(classification_report(Y_train,y_pred))

```

Local Outlier Factor: 291532

Accuracy Score :

0.053433726529194224

Classification Report :

	precision	recall	f1-score	support
0	0.03	0.89	0.05	9197
1	0.89	0.03	0.05	298792
accuracy			0.05	307989
macro avg	0.46	0.46	0.05	307989
weighted avg	0.87	0.05	0.05	307989

```
[ ]: n_outliers = len(Attack)
for i, (clf_name,clf) in enumerate(classifiers.items()):
    #Fit the data and tag outliers
    if clf_name == "Local Outlier Factor":
        continue
    y_pred = clf.fit_predict(X_train)
    scores_prediction = clf.negative_outlier_factor_
    elif clf_name == "Support Vector Machine":
        clf.fit(X_train)
        y_pred = clf.predict(X_train)
    else:
        continue
        clf.fit(X_train)
        scores_prediction = clf.decision_function(X_train)
        y_pred = clf.predict(X_train)

    #Reshape the prediction values to 0 for Valid transactions , 1 for Fraud
    ↪ transactions
    y_pred[y_pred == 1] = 0
    y_pred[y_pred == -1] = 1
    n_errors = (y_pred != Y_train).sum()
    # Run Classification Metrics
    print("{}: {}".format(clf_name,n_errors))
    print("Accuracy Score :")
    print(accuracy_score(Y_train,y_pred))
    print("Classification Report :")
    print(classification_report(Y_train,y_pred))
```

```
[12]: from keras.layers import Input, Dense
from keras.models import Model, Sequential
from keras import regularizers
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import classification_report, accuracy_score
from sklearn.manifold import TSNE
from sklearn import preprocessing
import matplotlib.pyplot as plt
import pandas as pd
import numpy as np
import seaborn as sns
```

```
[25]: def tsne_plot(x1, y1, name="graph.png"):

    #Scale features to improve the training ability of TSNE.
    standard_scaler = StandardScaler()
    df2_std = standard_scaler.fit_transform(x1)
```


```

tsne = TSNE(n_components=2, random_state=0)
x_test_2d = tsne.fit_transform(df2_std)

#Build the scatter plot with the two types of transactions.
color_map = {0:'green', 1:'red'}
plt.figure()
for idx, cl in enumerate(np.unique(y)):
    plt.scatter(x = x_test_2d[y1==cl,0],
                y = x_test_2d[y1==cl,1],
                c = color_map[idx],
                label = cl)
plt.xlabel('X in t-SNE')
plt.ylabel('Y in t-SNE')
plt.legend(loc='upper left')
plt.title('t-SNE visualization of test data')
plt.show()

tsne_plot(X_train[:8000], Y_train[:8000])

```



[27]: *#I think training this kind of dataset with autoencoders makes more sense so*  
*↳ let's try that*

```
## input layer
```

```

input_layer = Input(shape=(X_train.shape[1],))

## encoding part
encoded = Dense(100, activation='tanh', activity_regularizer=regularizers.
    ↳l1(10e-5))(input_layer)
encoded = Dense(50, activation='relu')(encoded)

## decoding part
decoded = Dense(50, activation='tanh')(encoded)
decoded = Dense(100, activation='tanh')(decoded)

## output layer
output_layer = Dense(X_train.shape[1], activation='relu')(decoded)

```

```

[28]: #Compile the mdoel
autoencoder = Model(input_layer, output_layer)
autoencoder.compile(optimizer="adadelta", loss="mse")

```

```

[29]: #let's do bit of data transformation for scaling

x = X_train
y = Y_train

x_scale = preprocessing.MinMaxScaler().fit_transform(x.values)
x_norm, x_fraud = x_scale[y == 0], x_scale[y == 1]

```

```

[30]: autoencoder.fit(x_norm[0:8000], x_norm[0:8000],
                    batch_size = 256, epochs = 50,
                    shuffle = True, validation_split = 0.20);

```

```

Epoch 1/50
25/25 [=====] - 0s 8ms/step - loss: 0.1274 - val_loss:
0.1273
Epoch 2/50
25/25 [=====] - 0s 4ms/step - loss: 0.1270 - val_loss:
0.1269
Epoch 3/50
25/25 [=====] - 0s 5ms/step - loss: 0.1266 - val_loss:
0.1264
Epoch 4/50
25/25 [=====] - 0s 5ms/step - loss: 0.1261 - val_loss:
0.1259
Epoch 5/50
25/25 [=====] - 0s 5ms/step - loss: 0.1255 - val_loss:
0.1253
Epoch 6/50
25/25 [=====] - 0s 6ms/step - loss: 0.1249 - val_loss:

```

0.1247  
Epoch 7/50  
25/25 [=====] - 0s 5ms/step - loss: 0.1243 - val\_loss:  
0.1240  
Epoch 8/50  
25/25 [=====] - 0s 6ms/step - loss: 0.1236 - val\_loss:  
0.1234  
Epoch 9/50  
25/25 [=====] - 0s 5ms/step - loss: 0.1230 - val\_loss:  
0.1228  
Epoch 10/50  
25/25 [=====] - 0s 11ms/step - loss: 0.1223 - val\_loss:  
0.1221  
Epoch 11/50  
25/25 [=====] - 0s 11ms/step - loss: 0.1216 - val\_loss:  
0.1214  
Epoch 12/50  
25/25 [=====] - 0s 7ms/step - loss: 0.1209 - val\_loss:  
0.1207  
Epoch 13/50  
25/25 [=====] - 0s 7ms/step - loss: 0.1202 - val\_loss:  
0.1199  
Epoch 14/50  
25/25 [=====] - 0s 7ms/step - loss: 0.1195 - val\_loss:  
0.1192  
Epoch 15/50  
25/25 [=====] - 0s 6ms/step - loss: 0.1188 - val\_loss:  
0.1185  
Epoch 16/50  
25/25 [=====] - 0s 8ms/step - loss: 0.1180 - val\_loss:  
0.1178  
Epoch 17/50  
25/25 [=====] - 0s 10ms/step - loss: 0.1173 - val\_loss:  
0.1171  
Epoch 18/50  
25/25 [=====] - 0s 8ms/step - loss: 0.1166 - val\_loss:  
0.1164  
Epoch 19/50  
25/25 [=====] - 0s 7ms/step - loss: 0.1159 - val\_loss:  
0.1157  
Epoch 20/50  
25/25 [=====] - 0s 7ms/step - loss: 0.1152 - val\_loss:  
0.1150  
Epoch 21/50  
25/25 [=====] - 0s 6ms/step - loss: 0.1144 - val\_loss:  
0.1142  
Epoch 22/50  
25/25 [=====] - 0s 8ms/step - loss: 0.1137 - val\_loss:



```

0.1135
Epoch 23/50
25/25 [=====] - 0s 10ms/step - loss: 0.1130 - val_loss:
0.1128
Epoch 24/50
25/25 [=====] - 0s 8ms/step - loss: 0.1123 - val_loss:
0.1122
Epoch 25/50
25/25 [=====] - 0s 7ms/step - loss: 0.1116 - val_loss:
0.1115
Epoch 26/50
25/25 [=====] - 0s 7ms/step - loss: 0.1109 - val_loss:
0.1108
Epoch 27/50
25/25 [=====] - 0s 7ms/step - loss: 0.1102 - val_loss:
0.1101
Epoch 28/50
25/25 [=====] - 0s 6ms/step - loss: 0.1096 - val_loss:
0.1094
Epoch 29/50
25/25 [=====] - 0s 5ms/step - loss: 0.1089 - val_loss:
0.1087
Epoch 30/50
25/25 [=====] - 0s 7ms/step - loss: 0.1082 - val_loss:
0.1081
Epoch 31/50
25/25 [=====] - 0s 7ms/step - loss: 0.1075 - val_loss:
0.1074
Epoch 32/50
25/25 [=====] - 0s 8ms/step - loss: 0.1069 - val_loss:
0.1068
Epoch 33/50
25/25 [=====] - 0s 9ms/step - loss: 0.1062 - val_loss:
0.1061
Epoch 34/50
25/25 [=====] - 0s 7ms/step - loss: 0.1056 - val_loss:
0.1055
Epoch 35/50
25/25 [=====] - 0s 7ms/step - loss: 0.1049 - val_loss:
0.1048
Epoch 36/50
25/25 [=====] - 0s 9ms/step - loss: 0.1043 - val_loss:
0.1042
Epoch 37/50
25/25 [=====] - 0s 10ms/step - loss: 0.1036 - val_loss:
0.1036
Epoch 38/50
25/25 [=====] - 0s 3ms/step - loss: 0.1030 - val_loss:

```

```

0.1029
Epoch 39/50
25/25 [=====] - 0s 3ms/step - loss: 0.1024 - val_loss:
0.1023
Epoch 40/50
25/25 [=====] - 0s 3ms/step - loss: 0.1017 - val_loss:
0.1017
Epoch 41/50
25/25 [=====] - 0s 3ms/step - loss: 0.1011 - val_loss:
0.1011
Epoch 42/50
25/25 [=====] - 0s 3ms/step - loss: 0.1005 - val_loss:
0.1004
Epoch 43/50
25/25 [=====] - 0s 7ms/step - loss: 0.0999 - val_loss:
0.0998
Epoch 44/50
25/25 [=====] - 0s 8ms/step - loss: 0.0992 - val_loss:
0.0992
Epoch 45/50
25/25 [=====] - 0s 10ms/step - loss: 0.0986 - val_loss:
0.0986
Epoch 46/50
25/25 [=====] - 0s 8ms/step - loss: 0.0980 - val_loss:
0.0980
Epoch 47/50
25/25 [=====] - 0s 8ms/step - loss: 0.0974 - val_loss:
0.0974
Epoch 48/50
25/25 [=====] - 0s 10ms/step - loss: 0.0968 - val_loss:
0.0968
Epoch 49/50
25/25 [=====] - 0s 8ms/step - loss: 0.0962 - val_loss:
0.0962
Epoch 50/50
25/25 [=====] - 0s 6ms/step - loss: 0.0956 - val_loss:
0.0956

```

[31]: *#Let's try to get latent learnt representation by autoencoder*

```

hidden_representation = Sequential()
hidden_representation.add(autoencoder.layers[0])
hidden_representation.add(autoencoder.layers[1])
hidden_representation.add(autoencoder.layers[2])

```

[32]:

```

norm_hid_rep = hidden_representation.predict(x_norm[9000:12000])
fraud_hid_rep = hidden_representation.predict(x_fraud[9000:12000])

```

```
[33]: #Let's try to visualize TSNE again for the learnt representation
```


```
rep_x = np.append(norm_hid_rep, fraud_hid_rep, axis = 0)
```

```
y_n = np.zeros(norm_hid_rep.shape[0])
```

```
y_f = np.ones(fraud_hid_rep.shape[0])
```

```
rep_y = np.append(y_n, y_f)
```

```
tsne_plot(rep_x, rep_y)
```



```
[34]: #Finally we can train a classifier on learnt representations
```

```
train_x, val_x, train_y, val_y = train_test_split(rep_x, rep_y, test_size=0.25)
```

```
clf = LogisticRegression(solver="lbfgs").fit(train_x, train_y)
```

```
pred_y = clf.predict(val_x)
```

```
print ("")
```

```
print ("Classification Report: ")
```

```
print (classification_report(val_y, pred_y))
```

```
print ("")
```

```
print ("Accuracy Score: ", accuracy_score(val_y, pred_y))
```

Classification Report:

precision	recall	f1-score	support
-----------	--------	----------	---------

	0.0	1.00	0.99	0.99	783
	1.0	0.99	1.00	0.99	717
accuracy				0.99	1500
macro avg		0.99	0.99	0.99	1500
weighted avg		0.99	0.99	0.99	1500

Accuracy Score: 0.9946666666666667

[37]: *#Finally we can train a classifier on learnt representations*

```
train_x, val_x, train_y, val_y = train_test_split(rep_x, rep_y, test_size=0.25)
clf = LogisticRegression(solver="lbfgs").fit(X_train, Y_train)
pred_y = clf.predict(X_val)

print ("")
print ("Classification Report: ")
print (classification_report(Y_val, pred_y))

print ("")
print ("Accuracy Score: ", accuracy_score(Y_val, pred_y))
```

/home/coding/MyProjects/Avira-AnomalyDetection/env/lib/python3.6/site-packages/sklearn/linear\_model/\_logistic.py:764: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

extra\_warning\_msg=\_LOGISTIC\_SOLVER\_CONVERGENCE\_MSG)

Classification Report:

	precision	recall	f1-score	support
0	1.00	0.99	0.99	82189
1	1.00	1.00	1.00	2689714
accuracy			1.00	2771903
macro avg	1.00	0.99	1.00	2771903
weighted avg	1.00	1.00	1.00	2771903

Accuracy Score: 0.9996976806186941

[ ]: